


What the harness assumes

concurrency = $\lceil C \cdot S \rceil$ workers, each turning a request round in exactly S

so capacity = concurrency / $S = C$

What a worker actually pays



true capacity = concurrency / ($S + \text{ov}$) = $C \cdot S / (S + \text{ov})$ — the same ov costs proportionally more when S is short