


What the harness assumes

concurrency = $\lceil C \cdot S \rceil$ workers, each turning a request round in exactly S

so capacity = concurrency / $S = C$

What a worker actually pays



$ov = 0.463$ ms,
unaccounted for in C

$C_{\text{model}} = \text{concurrency} / (S + ov) = C \cdot S / (S + ov)$ — the same ov costs proportionally more when S is short